

# CSCI 4160/6963, ECSE 4965/6965

## Reinforcement Learning

### Homework 9

#### Overview

This is a large-scale deep learning assignment. You will work in teams of 2 to 4. You will train a reasonably large convolutional neural network (CNN) to classify buildings on the RPI campus. We have collected images of 11 buildings on campus: 87gym, Amos Eaton, EMPAC, Greene, JEC, Lally, Library, Ricketts, Sage, Troy Building, Voorhees. There are roughly 500 images per building, though the numbers are not exactly the same. To make your life easier, I have shrunk the original images to a more manageable size,  $252 \times 189 \times 3$ , though you are welcome to resize them further before feeding them into the neural network. I have separated the data into training and test sets, and there is also a separate hidden test set that we will be using for grading.

To warm up, I have also provided a very small CNN for CIFAR10 that you can use as a starting point. You are required to reach a certain accuracy on CIFAR10 also. The challenges presented by CIFAR10 are by and large the same (modulo the size of the image), so understanding how to improve your CIFAR10 performance will greatly help for the big task.

#### Logistics

For this assignment, you will be using the NPL cluster at CCI. Since it is shared across RPI and beyond, CCI works by scheduling your jobs (i.e., programs). Every time you run a big-ish job, you should schedule it so that the machines don't get overwhelmed.

**Connecting to the cluster.** To connect to the cluster, you need to execute the `ssh` command (either from a Windows Powershell or a Unix command line, in the case of Mac/Linux). In all examples, I will use the generic account name, **RNL2user**. Your account name should start with the same first four letters (**RNL2** is the name of the project). Make sure to replace the generic account name with yours in all commands.

To connect, you first need to `ssh` to a landing pad. Landing pads are essentially CCI's interface with the outside world. These are not powerful machines and are not supposed to execute any heavy-duty tasks. Typically, I just `ssh` to a landing pad and then directly to a main machine. There are four landing pads: `blp01`, `blp02`, `blp03` and `blp04`. You can connect to either one as follows:

```
ssh RNL2user@blp01.ccni.rpi.edu
```

Once you enter the command, you will be asked for your CCI password, which you must set up online (this is not your RCS account, so you will need to set it up separately). Once you type your password, you will need to complete the two-factor authentication through Duo, which you also need to set up online beforehand.

Once you connect to a landing pad, you can connect to the main machines. There is only one option here, `nplfen01`:

```
ssh RNL2user@nplfen01.ccni.rpi.edu
```

This time you only need to enter your password. You can also set up a password-less `ssh` by following the instructions here: [https://docs.cci.rpi.edu/landingpads/Passwordless\\_SSH/](https://docs.cci.rpi.edu/landingpads/Passwordless_SSH/). Keep in mind that all machines mount the same file system, so you should see the same files no matter where you are (including the landing pads).

**CCI file system.** Your home directory has four folders: `barn`, `barn-shared`, `scratch`, `scratch-shared`. By convention, the data and installation files are in the `scratch` folders, and you code goes in the `barn` folders. **The `barn-shared` and `scratch-shared` folders are accessible by everyone, so be careful what you put there!**

Starting code is provided in the `~/barn-shared/startup_files` folder. Please do not modify anything there and just copy everything from that folder to your own `barn` folder.

Images are in the `~/scratch-shared/data` folder. You are welcome to copy images to your own `scratch` folder (or to your own machine if you wish to do small tests). Please do not distribute any images to people who are not in the class.

The conda environment is in the `~/scratch-shared/anaconda3` folder. You shouldn't have to touch this folder, except for the very first time you log on (see Getting started).

**Navigating the file system.** If you've never used a command line interface before, it might be hard initially. There are plenty of good resources online. The main commands you will need are `cp` (copy a file), `cd` (enter a directory), `ls` (list items in directory), `rm` (remove a file).

**Running your code.** Because CCI is a shared cluster, it works by scheduling jobs. While you can still run commands normally (e.g., `python train.py`), this should only be done for small tasks. If you run a bigger task this way, it may be killed and you are violating the responsible use agreement you signed. For bigger jobs, you need to use `srun`, which is the scheduling tool. The basic usage is as follows:

```
srun -t 1 --gres=gpu:1 python train_buildings.py
```

The `-t` flag indicates how many minutes your code can run for before it is killed. This can be conservative (e.g., a couple of hours). The GPU command gives you 1 GPU – you shouldn't have to change that. The rest is the actual command you would like to use.

**Moving files to and from the server.** To copy a file from your machine to CCI, use the `scp` command:

```
scp your_file.py RNL2user@blp01.ccni.rpi.edu:barn
```

You will once again need to do the 2-factor authentication. Note the folder after the colon (:); you always have to specify a folder when you `scp`.

To copy a file from CCI to your machine, use `sftp`:

```
sftp RNL2user@blp01.ccni.rpi.edu
```

After the 2-factor authentication, you are essentially `ssh'd` into `blp`, but you can't do much beyond moving around. If you want to copy a file to your machine, just call `get`:

```
get filename.py
```

The file will be copied to the directory where you called `sftp` from.

**Using command line text editors.** All command line text editors have a somewhat steep learning curve. The two most popular are `emacs` (my choice) and `vim`. Both have non-standard shortcuts and modes, if you have never used them before. Another choice is `nano`, which may be easier to use initially. You are also welcome to write code on your own machine and send it over, though constantly `scp`-ing may get annoying due to the 2-factor authentication.

**IMPORTANT: closing the command line interface.** Oftentimes, you may want to run your code, and then turn off your machine. That's possible, but you need to be careful. If you just close your command-line interface, your session will be interrupted and all progress will be lost. To avoid this, you should use `tmux`, which lets you start a session and detach from it without killing it. Using `tmux` is quite easy. Once you `ssh` to `nplfen`, you just type

```
tmux
```

This will put you in a new session. Note you will need to activate your conda environment again (see Getting started). After that, just run your code as normal (with `srun`). To detach from the session, just type `Ctrl+b` and then `d`. After this, you can safely close your terminal. If you want to get back to the session, `ssh` to `nplfen` again and type

```
tmux attach
```

This will bring you back. Note that `tmux` does not keep your entire print history, so keep that in mind. If you want to scroll up, type `Ctrl+b`, followed by `[`. To exit scrolling, press `q`. To exit the `tmux`, type `exit` (only do this once you no longer need the current session).

## Getting Started

To get started, you need to initialize your conda environment. This only needs to be done once, the first time you `ssh`. Navigate to the `~/scratch-shared/acandonda3/bin` folder and run the `conda init` command:

```
cd ~/scratch-shared/anaconda3/bin
./conda init
```

This will add conda to your path. Restart your connection after this.

Every time you wish to run your code, you need to activate the shared conda environment, which contains the pytorch installation, among other things. To activate it, just type (from any folder):

```
conda activate pytorch-env-shared
```

Now you're ready to run your training code. Keep in mind that the `train_buildings.py` script will crash initially because the CNN class is empty.

## Grading

The warm-up task is to get 70% test accuracy on CIFAR10. I have provided you with a very small CNN (1 convolutional and 1 fully-connected layer) that you can also use as an example to build your big CNN. Try to understand why your test accuracy doesn't increase after a certain point and how you can improve it. This will be worth a total of 10 points. **Your small CNN cannot have more than 4 hidden layers total [batch norm, dropout and max pool layers are not counted in the limit].**

The main task is to achieve 85% (90% for graduate students) test accuracy on the building dataset (as measured on the hidden test set). You will lose 5 percentage points for each point below the target. You are allowed to use any number of fully-connected and convolutional layers, as well as max pooling, batch normalization, dropout and weight decay. **No other layer types, such as residual blocks, are allowed.** You are also allowed to manipulate the training data in any way you see fit (you are **not allowed** to use complex external libraries such as `AutoAugment`) – I have normalized the data and converted it to pytorch floats to get you started (in fact, my normalization is borrowed from the CIFAR10, so it may not be perfect). **The only constraint is that your network cannot have more than 5 million parameters.**

The two best performing teams (on the hidden test set) will receive an extra 10 points.

Once you are done with your training, please provide short answers to the questions below:

1) How many parameters does your network have? Please provide your calculations, layer by layer. You can ignore batch normalization parameters – just count the convolution and fully connected ones.

2) How did you improve the CIFAR10 accuracy? What was the issue?

## Hints and Tips

- Start early. Training on the big dataset will be much slower than CIFAR10, since the images and networks will be bigger. Once you have set up a good architecture, you should expect to reach 85% within 25-30 epochs (or faster depending on your setup).
- Use `tmux`. It will make your life much easier when you need to work on other stuff.
- CCI is sometimes down and sometimes fully occupied. Do not leave stuff for the last day because it is very likely that CCI will not cooperate.
- Save your model often. I have provided a small example of how to save and load your code. You may want to keep track of how many epochs each model has been trained for.
- Start small. You might want to create a very small training set (e.g., 100 examples), to debug your architecture. Also use CIFAR10 as an example as much as possible.
- You may want to shrink the images somewhat. I have provided a resize factor to make your life easier. Keep in mind that ImageNet images, for example, are  $224 \times 224 \times 3$ , so the provided size can probably work fine as well. You can also plot your rescaled images to see what they look like.
- Check out the AlexNet (ImageNet) architecture online for an inspiration. Note that dimensions will not match. Also, you probably don't need huge fully-connected layers.
- Beware big fully connected layers. Chances are you will run out of memory (even on CCI!) the first time you switch from a convolutional to a fully connected layer. Always keep track of how many elements each layer has (and how many weights).
- Check out the `torchvision.transforms` for various ways to augment your training (but not testing!) data. Again, you are not allowed to use complex libraries we haven't covered in class, e.g., `AutoAugment`.
- You may want to have a small validation set to tune hyper-parameters, but keep in mind that the training set is not that big. You are probably fine tuning on the (provided) test set in this assignment.
- When you use `srun`, you should use set the `flush` flag in the print function; otherwise, your print statement may be buffered and will not show up in real time. For example, to print a variable `var`, call `print(var, flush=True)`.

## Submission

Please use LMS to submit a zip file containing 1) your training **.py** code, along with instructions on how to run it, 2) a small test **.py** that we can use on a given test set, 3) your trained model (in **.pth** file), and 4) a **.pdf** file containing your answers to the above questions. The deadline is **11:59pm, Tuesday, December 3**.

**Important:** 3 days after the deadline, I will ask you to put your code and trained models in the `barn-shared` folder, so that we don't have to move files back and forth when we grade.